

```
In [11]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import os
import seaborn as sns
from sklearn.metrics import accuracy_score, classification_report, confusion_m
from sklearn.ensemble import RandomForestClassifier
```

## Loading The Data

```
In [12]: df = pd.read_csv(r"D:\iris classification\Iris.csv")
df.head()
```

```
Out[12]:
```

|   | SepalLengthCm | SepalWidthCm | PetalLengthCm | PetalWidthCm | Species     |
|---|---------------|--------------|---------------|--------------|-------------|
| 0 | 5.1           | 3.5          | 1.4           | 0.2          | Iris-setosa |
| 1 | 4.9           | 3.0          | 1.4           | 0.2          | Iris-setosa |
| 2 | 4.7           | 3.2          | 1.3           | 0.2          | Iris-setosa |
| 3 | 4.6           | 3.1          | 1.5           | 0.2          | Iris-setosa |
| 4 | 5.0           | 3.6          | 1.4           | 0.2          | Iris-setosa |

```
In [13]: # data stats
df.describe()
```

```
Out[13]:
```

|       | SepalLengthCm | SepalWidthCm | PetalLengthCm | PetalWidthCm |
|-------|---------------|--------------|---------------|--------------|
| count | 150.000000    | 150.000000   | 150.000000    | 150.000000   |
| mean  | 5.843333      | 3.054000     | 3.758667      | 1.198667     |
| std   | 0.828066      | 0.433594     | 1.764420      | 0.763161     |
| min   | 4.300000      | 2.000000     | 1.000000      | 0.100000     |
| 25%   | 5.100000      | 2.800000     | 1.600000      | 0.300000     |
| 50%   | 5.800000      | 3.000000     | 4.350000      | 1.300000     |
| 75%   | 6.400000      | 3.300000     | 5.100000      | 1.800000     |
| max   | 7.900000      | 4.400000     | 6.900000      | 2.500000     |

```
In [14]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 5 columns):
 #   Column          Non-Null Count  Dtype  
---  -
 0   SepalLengthCm   150 non-null   float64
 1   SepalWidthCm    150 non-null   float64
 2   PetalLengthCm   150 non-null   float64
 3   PetalWidthCm    150 non-null   float64
 4   Species         150 non-null   object  
dtypes: float64(4), object(1)
memory usage: 6.0+ KB
```

```
In [15]: df['Species'].value_counts()
```

```
Out[15]: Species
Iris-setosa      50
Iris-versicolor  50
Iris-virginica   50
Name: count, dtype: int64
```

```
In [16]: df.shape
```

```
Out[16]: (150, 5)
```

## Data Pre-Processing


```
In [17]: # checking null value
df.isnull().sum()
```

```
Out[17]: SepalLengthCm    0
SepalWidthCm             0
PetalLengthCm            0
PetalWidthCm             0
Species                  0
dtype: int64
```

## Exploratory Analysis


```
In [18]: sns.histplot(df['SepalLengthCm'])  
plt.title('Sepal Length Histogram')  
plt.show()
```

C:\Users\lalit\anaconda3\lib\site-packages\seaborn\\_oldcore.py:1119: FutureWarning: use\_inf\_as\_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.  
with pd.option\_context('mode.use\_inf\_as\_na', True):



```
In [19]: sns.histplot(df['SepalWidthCm'])  
plt.title('Sepal Width Histogram')  
plt.show()
```


C:\Users\lalit\anaconda3\lib\site-packages\seaborn\\_oldcore.py:1119: FutureWarning: use\_inf\_as\_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.  
with pd.option\_context('mode.use\_inf\_as\_na', True):



```
In [20]: sns.histplot(df['PetalLengthCm'])  
plt.title('Petal Length Histogram')  
plt.show
```

C:\Users\lalit\anaconda3\lib\site-packages\seaborn\\_oldcore.py:1119: FutureWarning: use\_inf\_as\_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.  
with pd.option\_context('mode.use\_inf\_as\_na', True):

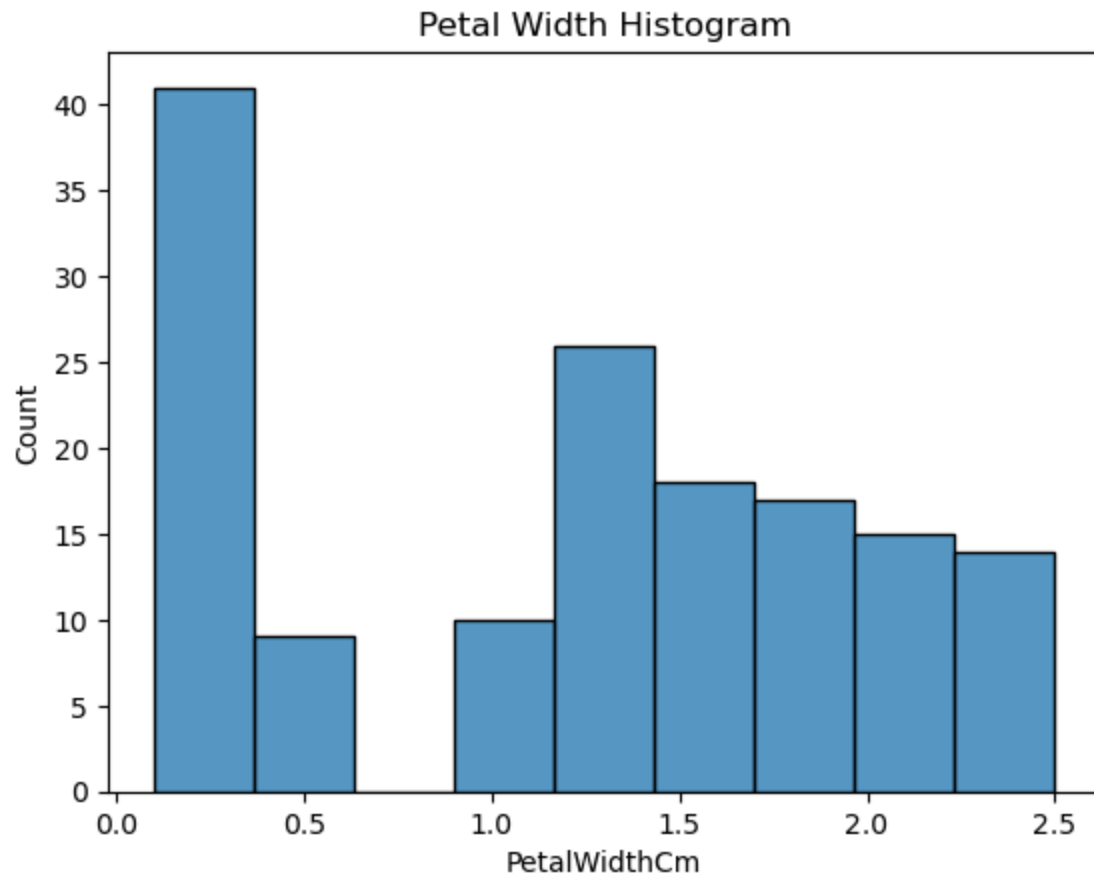
```
Out[20]: <function matplotlib.pyplot.show(close=None, block=None)>
```



```
In [21]: sns.histplot(df['PetalWidthCm'])  
plt.title('Petal Width Histogram')  
plt.show
```

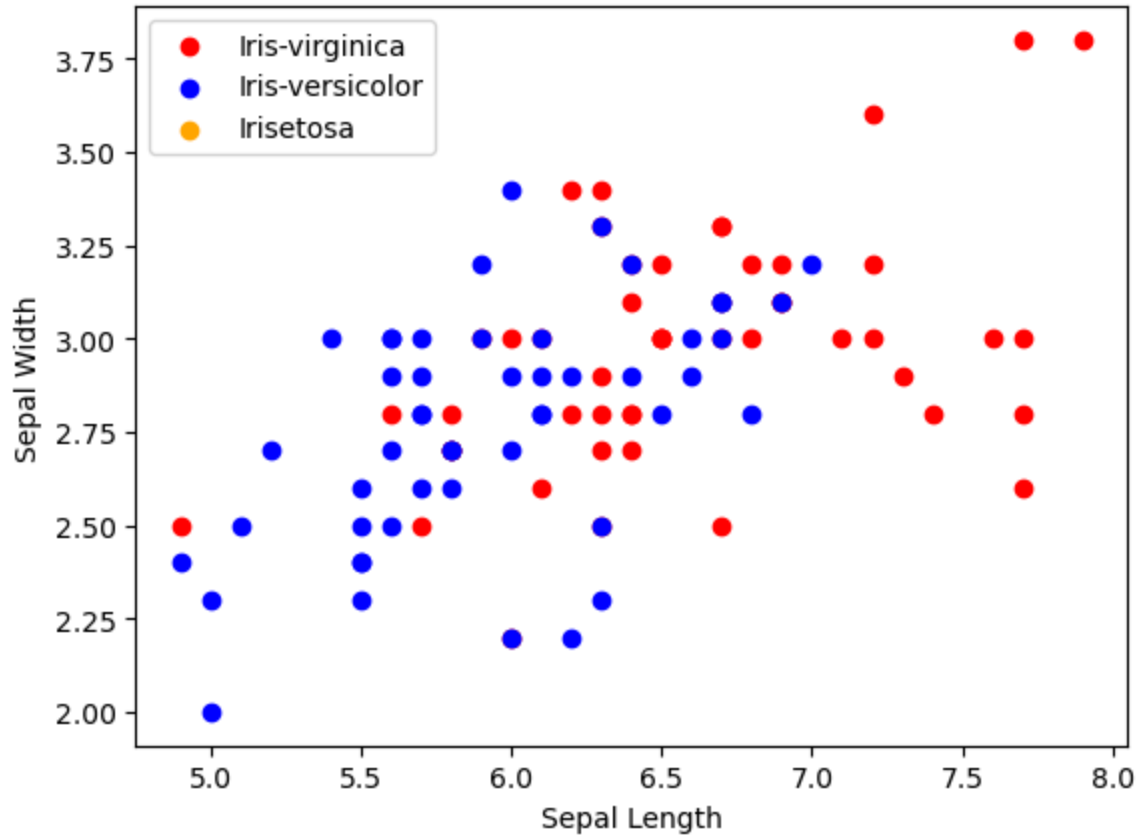
C:\Users\lalit\anaconda3\lib\site-packages\seaborn\\_oldcore.py:1119: FutureWarning: use\_inf\_as\_na option is deprecated and will be removed in a future version. Convert inf values to NaN before operating instead.  
with pd.option\_context('mode.use\_inf\_as\_na', True):

```
Out[21]: <function matplotlib.pyplot.show(close=None, block=None)>
```



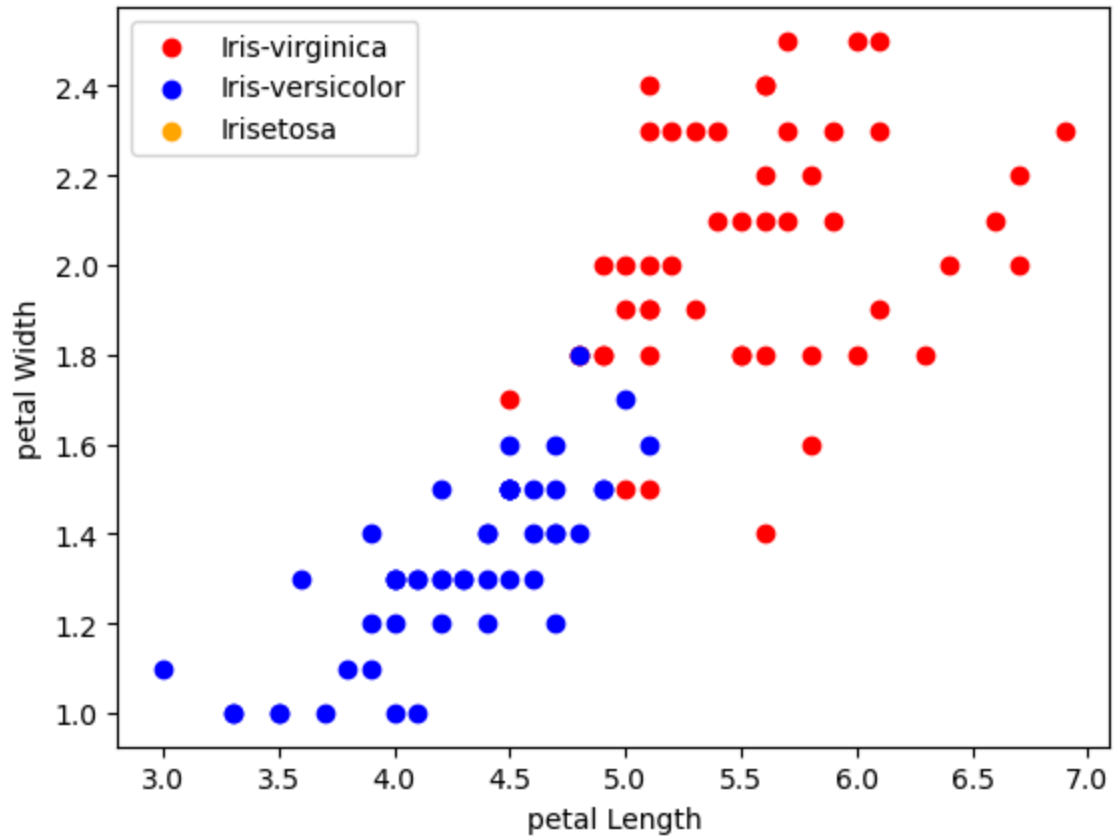
```
In [22]: # scatterplot  
color=['red','Blue','Orange']  
species=['Iris-virginica','Iris-versicolor','Irisetosa']
```

```
In [23]: for i in range(3):  
         x=df[df['Species']== species[i]]  
         plt.scatter(x['SepalLengthCm'],x['SepalWidthCm'],c=color[i],label=species[i])  
  
         plt.xlabel('Sepal Length')  
         plt.ylabel('Sepal Width')  
         plt.legend()
```



```
In [24]: for i in range(3):  
         x = df[df['Species'] == species[i]]  
         plt.scatter(x['PetalLengthCm'], x['PetalWidthCm'], c = color[i], label = s  
  
         plt.xlabel("petal Length")  
         plt.ylabel("petal Width")  
         plt.legend()
```

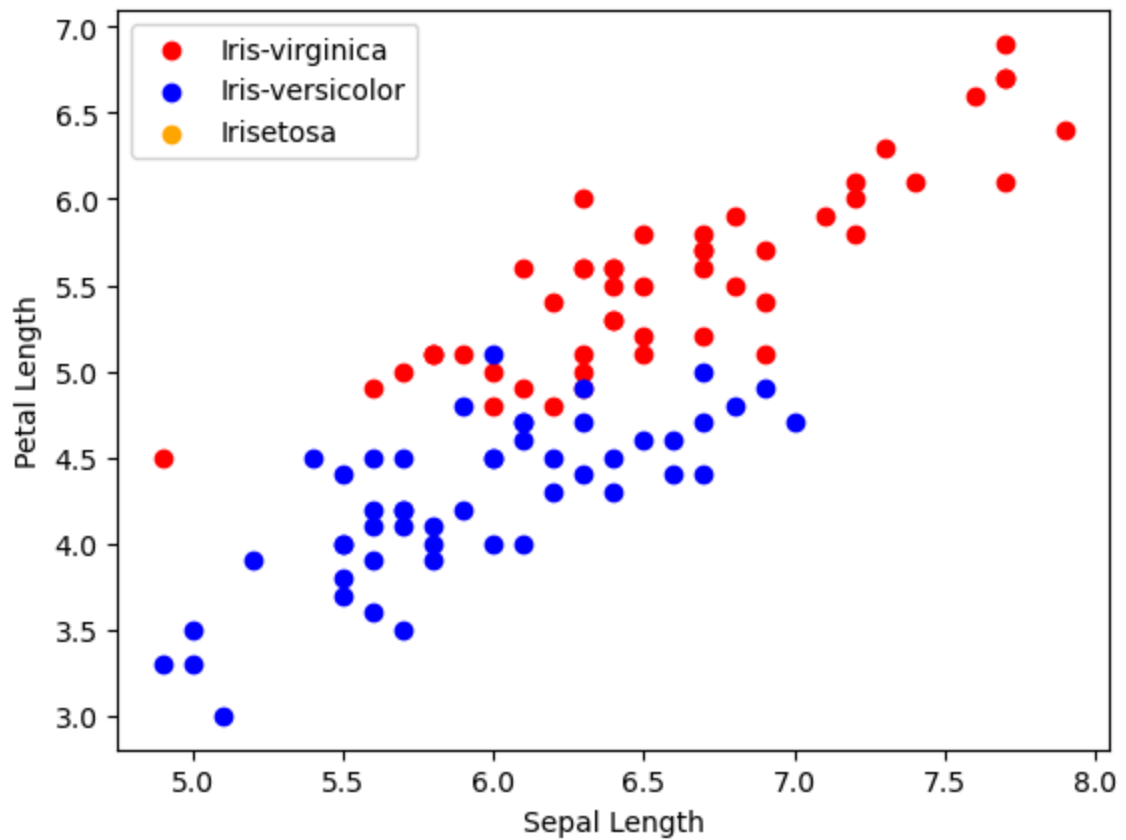
Out[24]: <matplotlib.legend.Legend at 0x22205091540>





```
In [25]: for i in range(3):  
         x = df[df['Species'] == species[i]]  
         plt.scatter(x['SepalLengthCm'], x['PetalLengthCm'], c = color[i], label =  
  
         plt.xlabel("Sepal Length")  
         plt.ylabel("Petal Length")  
         plt.legend()
```

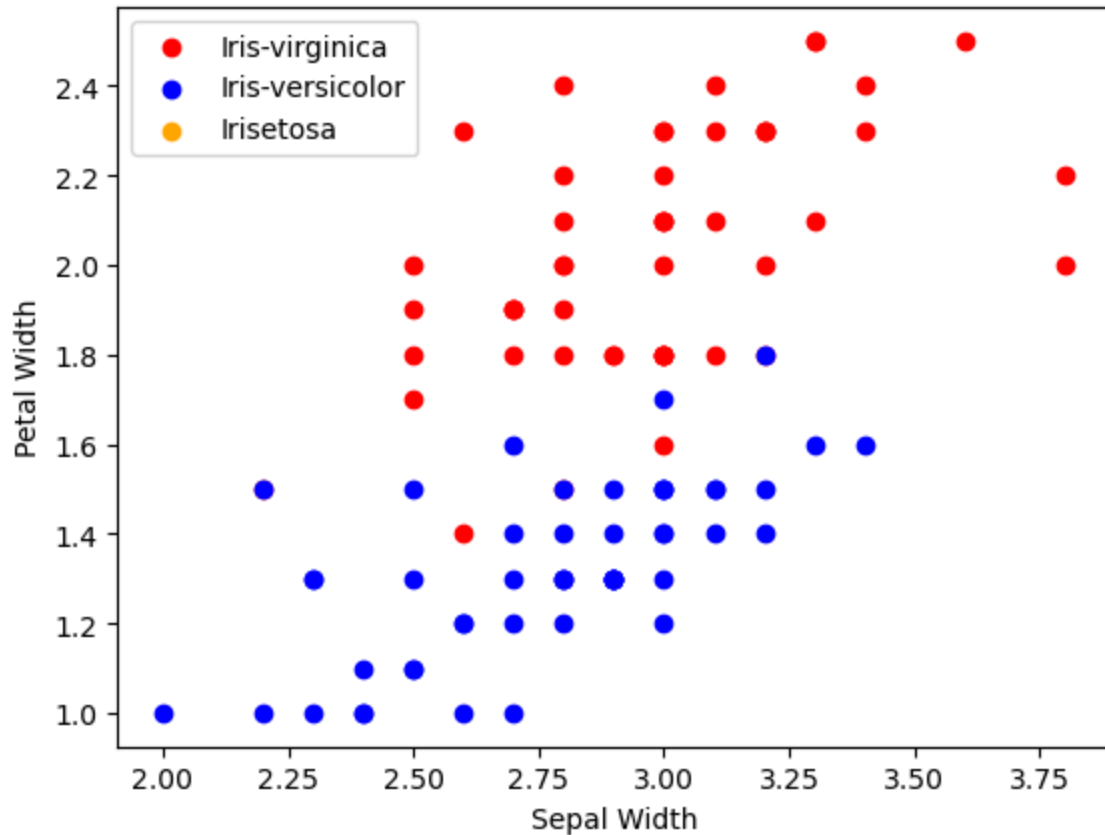
Out[25]: <matplotlib.legend.Legend at 0x22205337fd0>



```
In [26]: for i in range(3):
          x = df[df['Species'] == species[i]]
          plt.scatter(x['SepalWidthCm'], x['PetalWidthCm'], c = color[i], label = species[i])

          plt.xlabel("Sepal Width")
          plt.ylabel("Petal Width")
          plt.legend()
```

Out[26]: <matplotlib.legend.Legend at 0x2220536a8f0>



## Correlation Matrix


```
In [27]: numeric_df = df.select_dtypes(include=[np.number])
          numeric_df.corr()
```

Out[27]:

|               | SepalLengthCm | SepalWidthCm | PetalLengthCm | PetalWidthCm |
|---------------|---------------|--------------|---------------|--------------|
| SepalLengthCm | 1.000000      | -0.109369    | 0.871754      | 0.817954     |
| SepalWidthCm  | -0.109369     | 1.000000     | -0.420516     | -0.356544    |
| PetalLengthCm | 0.871754      | -0.420516    | 1.000000      | 0.962757     |
| PetalWidthCm  | 0.817954      | -0.356544    | 0.962757      | 1.000000     |

```
In [28]: corr = numeric_df.corr()
fig, ax = plt.subplots(figsize = (8,6))
sns.heatmap(corr, annot = True, ax = ax)
```

Out[28]: <Axes: >



## Label Encoder

```
In [29]: from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
df['Species'] = le.fit_transform(df['Species'])
```

```
In [30]: # Split data into features and target
X = df.drop(columns=["Species"])
Y = df['Species']
```

## Model Training

```
In [32]: from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(X, Y, test_size=0.30, rand
```

```
In [34]: #random Forest
rf_model = RandomForestClassifier(n_estimators=100, random_state=42)
rf_model.fit(x_train, y_train)

y_pred_rf = rf_model.predict(x_test)
accuracy_rf = accuracy_score(y_test, y_pred_rf)
print(f"Random Forest Accuracy: {accuracy_rf * 100:.2f}%")
print("Classification Report:", classification_report(y_test, y_pred_rf))
print("Confusion Matrix:", confusion_matrix(y_test, y_pred_rf))
```

Random Forest Accuracy: 88.89%

| Classification Report: | precision | recall | f1-score | support |
|------------------------|-----------|--------|----------|---------|
| 0                      | 1.00      | 1.00   | 1.00     | 15      |
| 1                      | 0.78      | 0.93   | 0.85     | 15      |
| 2                      | 0.92      | 0.73   | 0.81     | 15      |
| accuracy               |           | 0.89   |          | 45      |
| macro avg              | 0.90      | 0.89   | 0.89     | 45      |
| weighted avg           | 0.90      | 0.89   | 0.89     | 45      |

Confusion Matrix: [[15 0 0]  
[ 0 14 1]  
[ 0 4 11]]

```

In [35]: #hyper parameter tuned random forest
from sklearn.model_selection import GridSearchCV

param_grid = {
    'n_estimators': [100, 200, 300],
    'max_depth': [10, 20, 30],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'bootstrap': [True, False]
}

rf = RandomForestClassifier(random_state=42)

grid_search = GridSearchCV(estimator=rf, param_grid=param_grid, cv=5, n_jobs=-1)

grid_search.fit(x_train, y_train)

print("Best Hyperparameters:", grid_search.best_params_)

best_rf = grid_search.best_estimator_

y_pred = best_rf.predict(x_test)
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy with best parameters: {accuracy * 100:.2f}%")

print("Classification Report:")
print(classification_report(y_test, y_pred))
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))

```

Fitting 5 folds for each of 162 candidates, totalling 810 fits

Best Hyperparameters: {'bootstrap': True, 'max\_depth': 10, 'min\_samples\_leaf': 4, 'min\_samples\_split': 2, 'n\_estimators': 100}

Accuracy with best parameters: 91.11%

Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 1.00      | 1.00   | 1.00     | 15      |
| 1            | 0.82      | 0.93   | 0.87     | 15      |
| 2            | 0.92      | 0.80   | 0.86     | 15      |
| accuracy     |           |        | 0.91     | 45      |
| macro avg    | 0.92      | 0.91   | 0.91     | 45      |
| weighted avg | 0.92      | 0.91   | 0.91     | 45      |

Confusion Matrix:

```

[[15  0  0]
 [ 0 14  1]
 [ 0  3 12]]

```

```
In [36]: # logistic regression
from sklearn.linear_model import LogisticRegression
model = LogisticRegression(max_iter=200)
model.fit(x_train, y_train)

y_pred = model.predict(x_test)
accuracy = accuracy_score(y_test, y_pred)

print(f"Accuracy: {accuracy * 100:.2f}%")

print("Classification Report:")
print(classification_report(y_test, y_pred))

print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))
```

Accuracy: 93.33%

Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 1.00      | 1.00   | 1.00     | 15      |
| 1            | 0.88      | 0.93   | 0.90     | 15      |
| 2            | 0.93      | 0.87   | 0.90     | 15      |
| accuracy     |           |        | 0.93     | 45      |
| macro avg    | 0.93      | 0.93   | 0.93     | 45      |
| weighted avg | 0.93      | 0.93   | 0.93     | 45      |

Confusion Matrix:

```
[[15  0  0]
 [ 0 14  1]
 [ 0  2 13]]
```

```
In [37]: # hyperparameter tuned logistic regression
param_grid = {
    'C': [0.01, 0.1, 1, 10, 100],
    'penalty': ['l1', 'l2', 'elasticnet', 'none'],
    'solver': ['liblinear', 'saga'],
    'max_iter': [100, 200, 500],
}

log_reg = LogisticRegression()

grid_search = GridSearchCV(estimator=log_reg, param_grid=param_grid, cv=5, n_j

grid_search.fit(x_train, y_train)

print("Best Hyperparameters:", grid_search.best_params_)

best_log_reg = grid_search.best_estimator_

print("Classification Report:")
print(classification_report(y_test, y_pred))
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))
```

Fitting 5 folds for each of 120 candidates, totalling 600 fits

Best Hyperparameters: {'C': 0.01, 'max\_iter': 100, 'penalty': 'none', 'solver': 'saga'}

Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 1.00      | 1.00   | 1.00     | 15      |
| 1            | 0.88      | 0.93   | 0.90     | 15      |
| 2            | 0.93      | 0.87   | 0.90     | 15      |
| accuracy     |           |        | 0.93     | 45      |
| macro avg    | 0.93      | 0.93   | 0.93     | 45      |
| weighted avg | 0.93      | 0.93   | 0.93     | 45      |

Confusion Matrix:

```
[[15  0  0]
 [ 0 14  1]
 [ 0  2 13]]
```

C:\Users\lalit\anaconda3\lib\site-packages\sklearn\model\_selection\\_valida

```
In [38]: # knn - K-nearest neighbor
from sklearn.neighbors import KNeighborsClassifier
model = KNeighborsClassifier()
model.fit(x_train, y_train)

y_pred = model.predict(x_test)

accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy*100}")

print("Classification Report:")
print(classification_report(y_test, y_pred))

print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))
```

Accuracy: 97.77777777777777

Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 1.00      | 1.00   | 1.00     | 15      |
| 1            | 0.94      | 1.00   | 0.97     | 15      |
| 2            | 1.00      | 0.93   | 0.97     | 15      |
| accuracy     |           |        | 0.98     | 45      |
| macro avg    | 0.98      | 0.98   | 0.98     | 45      |
| weighted avg | 0.98      | 0.98   | 0.98     | 45      |

Confusion Matrix:

```
[[15  0  0]
 [ 0 15  0]
 [ 0  1 14]]
```



```

In [39]: # hyperparameter tuned KNN
param_grid = {
    'n_neighbors': [3, 5, 7, 9],
    'weights': ['uniform', 'distance'],
    'metric': ['euclidean', 'manhattan', 'minkowski']
}

knn = KNeighborsClassifier()

grid_search = GridSearchCV(estimator=knn, param_grid=param_grid, cv=5, n_jobs=
grid_search.fit(x_train, y_train)

print("Best Hyperparameters:", grid_search.best_params_)

best_knn = grid_search

print("Classification Report:")
print(classification_report(y_test, y_pred))
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))

```

Fitting 5 folds for each of 24 candidates, totalling 120 fits

Best Hyperparameters: {'metric': 'euclidean', 'n\_neighbors': 9, 'weights': 'uniform'}

Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 1.00      | 1.00   | 1.00     | 15      |
| 1            | 0.94      | 1.00   | 0.97     | 15      |
| 2            | 1.00      | 0.93   | 0.97     | 15      |
| accuracy     |           |        | 0.98     | 45      |
| macro avg    | 0.98      | 0.98   | 0.98     | 45      |
| weighted avg | 0.98      | 0.98   | 0.98     | 45      |

Confusion Matrix:

```

[[15  0  0]
 [ 0 15  0]
 [ 0  1 14]]

```

```
In [43]: # decision tree
from sklearn.tree import DecisionTreeClassifier
model = DecisionTreeClassifier()
model.fit(x_train, y_train)
y_pred = model.predict(x_test)

accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy*100}")

print("Classification Report:")
print(classification_report(y_test, y_pred))

print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))
```

Accuracy: 93.33333333333333

Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 1.00      | 1.00   | 1.00     | 15      |
| 1            | 0.88      | 0.93   | 0.90     | 15      |
| 2            | 0.93      | 0.87   | 0.90     | 15      |
| accuracy     |           |        | 0.93     | 45      |
| macro avg    | 0.93      | 0.93   | 0.93     | 45      |
| weighted avg | 0.93      | 0.93   | 0.93     | 45      |

Confusion Matrix:

```
[[15  0  0]
 [ 0 14  1]
 [ 0  2 13]]
```

```
In [44]: #hyperparameter tuned decision tree
param_grid = {
    'criterion': ['gini', 'entropy'],
    'max_depth': [5, 10, 20, None],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'max_features': [None, 'auto', 'sqrt', 'log2']
}

dt = DecisionTreeClassifier(random_state=42)

grid_search = GridSearchCV(estimator=dt, param_grid=param_grid, cv=5, n_jobs=-1)
grid_search.fit(x_train, y_train)

print("Best Hyperparameters:", grid_search.best_params_)

best_dt = grid_search.best_estimator_

y_pred = best_dt.predict(x_test)
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy with best parameters: {accuracy * 100:.2f}%")

print("Classification Report:")
print(classification_report(y_test, y_pred))
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))
```

Fitting 5 folds for each of 288 candidates, totalling 1440 fits

Best Hyperparameters: {'criterion': 'gini', 'max\_depth': 5, 'max\_features': 'auto', 'min\_samples\_leaf': 1, 'min\_samples\_split': 5}

Accuracy with best parameters: 84.44%

Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 1.00      | 1.00   | 1.00     | 15      |
| 1            | 0.72      | 0.87   | 0.79     | 15      |
| 2            | 0.83      | 0.67   | 0.74     | 15      |
| accuracy     |           |        | 0.84     | 45      |
| macro avg    | 0.85      | 0.84   | 0.84     | 45      |
| weighted avg | 0.85      | 0.84   | 0.84     | 45      |

Confusion Matrix:

```
[[15  0  0]
 [ 0 13  2]
 [ 0  5 10]]
```

C:\Users\lalit\anaconda3\lib\site-packages\sklearn\tree\\_classes.py:269: FutureWarning: `max\_features='auto'` has been deprecated in 1.1 and will be removed in 1.3. To keep the past behaviour, explicitly set `max\_features='sqrt'`.

```
warnings.warn(
```

```

In [45]: #SVM
from sklearn.svm import SVC

model = SVC(kernel='linear', random_state=42)
model.fit(x_train, y_train)

y_pred = model.predict(x_test)

accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy*100}")

print("Classification Report:")
print(classification_report(y_test, y_pred))

print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))

```

Accuracy: 100.0

Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 1.00      | 1.00   | 1.00     | 15      |
| 1            | 1.00      | 1.00   | 1.00     | 15      |
| 2            | 1.00      | 1.00   | 1.00     | 15      |
| accuracy     |           |        | 1.00     | 45      |
| macro avg    | 1.00      | 1.00   | 1.00     | 45      |
| weighted avg | 1.00      | 1.00   | 1.00     | 45      |

Confusion Matrix:

```

[[15  0  0]
 [ 0 15  0]
 [ 0  0 15]]

```

```

In [46]: #hyperparameter tuned SVM
param_grid = {
    'C': [0.1, 1, 10, 100],
    'kernel': ['linear', 'rbf', 'poly'],
    'gamma': ['scale', 'auto'],
    'degree': [2, 3, 4],
}
svc = SVC()

grid_search = GridSearchCV(estimator=svc, param_grid=param_grid, cv=5, verbose=1)
grid_search.fit(x_train, y_train)
print("Best Hyperparameters:", grid_search.best_params_)
best_svm = grid_search.best_estimator_

y_pred = best_svm.predict(x_test)

accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy with best parameters: {accuracy * 100:.2f}%")
print("Classification Report:")
print(classification_report(y_test, y_pred))
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))

```

Fitting 5 folds for each of 72 candidates, totalling 360 fits

Best Hyperparameters: {'C': 0.1, 'degree': 2, 'gamma': 'scale', 'kernel': 'linear'}

Accuracy with best parameters: 95.56%

Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 1.00      | 1.00   | 1.00     | 15      |
| 1            | 0.93      | 0.93   | 0.93     | 15      |
| 2            | 0.93      | 0.93   | 0.93     | 15      |
| accuracy     |           |        | 0.96     | 45      |
| macro avg    | 0.96      | 0.96   | 0.96     | 45      |
| weighted avg | 0.96      | 0.96   | 0.96     | 45      |

Confusion Matrix:

```

[[15  0  0]
 [ 0 14  1]
 [ 0  1 14]]

```

In [ ]: